

Application of learning algorithms to nonlinear filtering and MCMC

PhD Dissertation defense

Nov 6, 2019

Anand Radhakrishnan



Department of Electrical and Computer Engineering — University of Florida

Committee chair: Prof. Sean Meyn

Committee members : Prof. José Principe, Prof. Kamran Mohseni, Prof. James Hobert

Thanks to the National Science Foundation and Army Research Office

Thanks to friends & family

List of Publications

Feedback Particle Filter:

- A. Radhakrishnan, A. M. Devraj, S. P. Meyn, **"Learning Techniques for Feedback Particle Filter Design"** *IEEE Conference on Decision and Control, Dec 2016.*
- A. Radhakrishnan, and S. P. Meyn, **"Feedback Particle Filter Design Using a Differential-Loss Reproducing Kernel Hilbert Space"** *American Control Conference, June 2018.*
- A. Radhakrishnan, and S. P. Meyn, **"Gain Function Tracking in the Feedback Particle Filter"** *American Control Conference, July 2019.*

Markov chain Monte Carlo methods:

- N. Brosse, A. Durmus, S. P. Meyn, E. Moulines, and A. Radhakrishnan, **"Diffusion Approximation and Control Variates for MCMC"** Submitted to *Annals of Applied Probability* in July 2019.

Code on Github

- FPF package in Julia by S. Surace - <https://github.com/simsurace/FeedbackParticleFilters.jl>
- Matlab and Python code - <https://github.com/a4anandr/FPF-code>

Outline

- 1 PhD Proposal - Recap
- 2 Poisson's Equation
- 3 Differential LSTD Learning
- 4 Applications to Nonlinear Filtering
- 5 Applications to MCMC
- 6 Conclusions and Future Work

PhD Proposal - Recap

Work till then

- Presented the **feedback particle filter (FPF)** - an approximation of the nonlinear filter that requires computing a gain function by solving a version of *Poisson's equation*.
- Proposed a **differential LSTD-learning** algorithm to approximate the gradient of the solution. [R et al. 16]

PhD Proposal - Recap

Work till then

- Presented the **feedback particle filter (FPF)** - an approximation of the nonlinear filter that requires computing a gain function by solving a version of *Poisson's equation*.
- Proposed a **differential LSTD-learning** algorithm to approximate the gradient of the solution. [R et al. 16]
- Presented numerical examples for gain approximation and state estimation for scalar systems.

PhD Proposal - Recap

Tasks promised and accomplished

- Extend the algorithm to a multidimensional setting

PhD Proposal - Recap

Tasks promised and accomplished

- Extend the algorithm to a multidimensional setting ✓
- Explore appropriate basis selection for these algorithms

PhD Proposal - Recap

Tasks promised and accomplished

- Extend the algorithm to a multidimensional setting ✓
 - Explore appropriate basis selection for these algorithms ✓
 - Impact of gain approximation error on filtering performance
- } [R et al. 18]

PhD Proposal - Recap

Tasks promised and accomplished

- Extend the algorithm to a multidimensional setting ✓
 - Explore appropriate basis selection for these algorithms ✓
 - Impact of gain approximation error on filtering performance - ✗
- [R et al. 18]
- [Taghvaei et al. 18]

PhD Proposal - Recap

Tasks promised and accomplished

- Extend the algorithm to a multidimensional setting ✓
 - Explore appropriate basis selection for these algorithms ✓
 - Impact of gain approximation error on filtering performance - X
- [R et al. 18]
- [Taghvaei et al. 18]

New directions

- Simplified differential LSTD algorithm for Langevin diffusion
- Application of kernel methods to gain approximation [R et al. 18]
- Improvements for online gain estimation [R et al. 19]
- Explored the application to MCMC algorithms [Brosse et al. 19]

Poisson's Equation

$$0 = \tilde{c} + \mathcal{D}h$$

$$h(x) = \mathbb{E} \left[\int_0^\tau \tilde{c}(X(t)) dt \right]$$

with $X(0) = x$

Optimal FPF Gain

$$K = \nabla h$$

$$\left\{ (x)^{u_0} \mathcal{L} + (n^+ x)^{\frac{n}{n-1}} \left\{ \phi(x) \right\}^{\frac{n}{n-1}} \right\} = (x)^{1+u_0} \phi$$

Optimal Control

$$\|z\|_{\theta}^2 \Delta \|z\|_{\theta}^2 =$$

$$\langle \theta^2, \theta^4 \rangle_{\theta} = \frac{\theta^2}{2}$$

$$\langle \theta^2, \theta^4 \Delta^2 h \rangle = \frac{1}{2} \mathcal{L}^2 \mathcal{L}$$

Optimal MCMC CV

Poisson's Equation

Poisson's Equation

- Second order partial differential equation with applications in various fields.
- General form in stochastic systems

$$\mathcal{D}h = -f$$

\mathcal{D} - differential operator

f - forcing function, usually centered

Poisson's Equation

- Second order partial differential equation with applications in various fields.
- General form in stochastic systems

$$\mathcal{D}h = -f$$

\mathcal{D} - differential operator

f - forcing function, usually centered

h - solution to Poisson's equation

Poisson's Equation

Stochastic optimal control

Example: In average cost optimal control problems,

$$c(x, u) + P_u h(x) = h(x) + \eta$$

$c(x, u)$ - Cost function associated with state x and action u .

P_u - Transition kernel of the controlled Markov chain.

η - Average cost.

h - **Relative value function**

Poisson's equation is the average-cost dynamic programming equation.

Poisson's Equation

Langevin Diffusion

Langevin diffusion is given by the SDE,

$$d\Phi_t = \underbrace{-\nabla U(\Phi_t) dt}_{\text{Drift term}} + \underbrace{\sqrt{2} dW_t}_{\text{Diffusion term}}, \quad \Phi \in \mathbb{R}^d$$

$U \in C^1$ is called the *potential function*.

$W = \{W_t : t \geq 0\}$ is a standard Brownian motion on \mathbb{R}^d .

Poisson's Equation

Langevin Diffusion

Langevin diffusion is given by the SDE,

$$d\Phi_t = \underbrace{-\nabla U(\Phi_t) dt}_{\text{Drift term}} + \underbrace{\sqrt{2} dW_t}_{\text{Diffusion term}}, \quad \Phi \in \mathbb{R}^d$$

$U \in C^1$ is called the *potential function*.

$W = \{W_t : t \geq 0\}$ is a standard Brownian motion on \mathbb{R}^d .

- May be regarded as a d -dimensional gradient flow with “noise”.

Poisson's Equation

Langevin Diffusion

Langevin diffusion is given by the SDE,

$$d\Phi_t = \underbrace{-\nabla U(\Phi_t) dt}_{\text{Drift term}} + \underbrace{\sqrt{2} dW_t}_{\text{Diffusion term}}, \quad \Phi \in \mathbb{R}^d$$

$U \in C^1$ is called the *potential function*.

$W = \{W_t : t \geq 0\}$ is a standard Brownian motion on \mathbb{R}^d .

- May be regarded as a d -dimensional gradient flow with “noise”.
- Diffusion is reversible, with unique *invariant density* $\rho = e^{-U+\Lambda}$, where Λ is a normalizing constant.

Poisson's Equation

Langevin Diffusion

Differential generator:

$$\mathcal{D}f := \lim_{t \rightarrow 0} \frac{\mathbb{E}[f(\Phi_t) | \Phi_0 = x] - f(x)}{t}$$

Poisson's Equation

Langevin Diffusion

Differential generator:

$$\begin{aligned}\mathcal{D}f &:= \lim_{t \rightarrow 0} \frac{\mathbb{E}[f(\Phi_t) | \Phi_0 = x] - f(x)}{t} \\ &= -\nabla U \cdot \nabla f + \Delta f, \quad f \in C^2,\end{aligned}$$

where ∇ is the gradient and Δ is the Laplacian.

Poisson's Equation

Langevin Diffusion

Differential generator:

$$\begin{aligned}\mathcal{D}f &:= \lim_{t \rightarrow 0} \frac{\mathbb{E}[f(\Phi_t) | \Phi_0 = x] - f(x)}{t} \\ &= -\nabla U \cdot \nabla f + \Delta f, \quad f \in C^2,\end{aligned}$$

where ∇ is the gradient and Δ is the Laplacian.

Let $c: \mathbb{R}^d \rightarrow \mathbb{R}$ be a function of interest, and

$$\eta = \int c(x) \rho(x) dx = \langle c, 1 \rangle_{L^2}.$$

Poisson's Equation

Langevin Diffusion

Differential generator:

$$\begin{aligned}\mathcal{D}f &:= \lim_{t \rightarrow 0} \frac{\mathbb{E}[f(\Phi_t) | \Phi_0 = x] - f(x)}{t} \\ &= -\nabla U \cdot \nabla f + \Delta f, \quad f \in C^2,\end{aligned}$$

where ∇ is the gradient and Δ is the Laplacian.

Let $c: \mathbb{R}^d \rightarrow \mathbb{R}$ be a function of interest, and

$$\eta = \int c(x) \rho(x) dx = \langle c, 1 \rangle_{L^2}.$$

Function $h \in C^2$ solves **Poisson's equation** with forcing function c if

$$\mathcal{D}h := -\tilde{c}, \quad \tilde{c} = c - \eta.$$

$$h := \int_0^\infty \mathbb{E}[\tilde{c}(\Phi_t)] dt$$

Poisson's Equation

Existence of a solution

- A solution exists under weak assumptions on U and c [Glynn & Meyn 96, Kontoyiannis et al. 12].
- Representations for the gradient of h and bounds are obtained in [Laugesen et al. 15, Devraj et al. 18].
- A smooth solution $h \in C^2$ exists under stronger conditions in [Pardoux et al. 01], subject to growth conditions on c similar to [Glynn & Meyn 96].

Poisson's Equation

Approximate solution to Poisson's equation

Obtaining an analytical solution for h is difficult outside special cases.
Hence approximation.

Poisson's Equation

Approximate solution to Poisson's equation

Obtaining an analytical solution for h is difficult outside special cases.
Hence approximation.

Goal: For a given function class \mathcal{H} , find the minimizer of

$$g^* := \arg \min_{g \in \mathcal{H}} \|h - g\|_{L^2}^2 \quad (\star)$$

Such minimum norm optimization problems can be solved using *TD learning* [Tsitsiklis 99].

Poisson's Equation

Approximate solution to Poisson's equation

Obtaining an analytical solution for h is difficult outside special cases.
Hence approximation.

Goal: For a given function class \mathcal{H} , find the minimizer of

$$g^* := \arg \min_{g \in \mathcal{H}} \|h - g\|_{L^2}^2 \quad (\star)$$

Such minimum norm optimization problems can be solved using *TD learning* [Tsitsiklis 99].

Challenge - No algorithm exists to solve (\star) if the process does not regenerate (for diffusions, $\dim > 1$ ruled out).

LSTD Learning

Discounted cost case

Discounted-cost value function:

$$h^\gamma(x) := \int_0^\infty e^{-\gamma t} \mathbb{E}_x[c(\Phi_t)] dt, \quad \gamma > 0 : \text{discount factor}$$

Discounted-cost optimality equation:

$$\gamma h^\gamma = c + \mathcal{D}h^\gamma$$

LSTD Learning

Discounted cost case

Discounted-cost value function:

$$h^\gamma(x) := \int_0^\infty e^{-\gamma t} \mathbb{E}_x[c(\Phi_t)] dt, \quad \gamma > 0 : \text{discount factor}$$

Discounted-cost optimality equation:

$$\gamma h^\gamma = c + \mathcal{D}h^\gamma$$

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \psi_i, \psi_j \rangle_{L^2}, \quad b_i = \langle \psi_i, h^\gamma \rangle_{L^2}$$

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \psi_i, \psi_j \rangle_{L^2}, \quad b_i = \langle \psi_i, h^\gamma \rangle_{L^2}$$

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$\begin{aligned} M_{ij} &= \langle \psi_i, \psi_j \rangle_{L^2}, & b_i &= \langle \psi_i, h^\gamma \rangle_{L^2} \\ & & &= \langle \psi_i, R_\gamma c \rangle_{L^2} \end{aligned}$$

Resolvent kernel: $R_\gamma c(x) := \int_0^\infty \mathbb{E}_x \left[e^{-\gamma t} c(\Phi_t) \right] dt$

$$R_\gamma c = (I\gamma - \mathcal{D})^{-1}c$$

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \psi_i, \psi_j \rangle_{L^2}, \quad b_i = \langle \psi_i, R_\gamma c \rangle_{L^2}$$

$$= \langle R_\gamma^\dagger \psi_i, c \rangle_{L^2}$$

Using an adjoint operation and applying the stationarity of Φ .

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \psi_i, \psi_j \rangle_{L^2}, \quad b_i = \langle R_\gamma^\dagger \psi_i, c \rangle_{L^2}$$

Eligibility vector: $\varphi(t) := \int_0^\infty e^{-\gamma r} \psi(\Phi_{t-r}) \mathrm{d}r$

$$R_\gamma^\dagger \psi_i(x) = \mathbb{E}[\varphi_i(t) | \Phi_t = x]$$

LSTD Learning

Discounted cost case

LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|h^\gamma - g\|_{L^2}^2$

For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i$$

$$\theta^* = M^{-1}b$$

$$\begin{aligned} M_{ij} &= \langle \psi_i, \psi_j \rangle_{L^2}, & b_i &= \langle R_\gamma^\dagger \psi_i, c \rangle_{L^2} \\ & & &= \mathbb{E}[\varphi_i(t)c(\Phi_t)] \end{aligned}$$

LSTD Learning

Discounted cost case

ODE formulation of the LSTD algorithm:

$$\frac{d}{dt}M(t) = \psi(\Phi_t)\psi^T(\Phi_t)$$

$$\frac{d}{dt}\varphi(t) = -\gamma\varphi(t) + \psi(\Phi_t)$$

$$\frac{d}{dt}b(t) = \varphi(t)c(\Phi_t)$$

$$\theta(t) := M(t)^{-1}b(t)$$

LSTD Learning

Discounted cost case

ODE formulation of the LSTD algorithm:

$$\frac{d}{dt}M(t) = \psi(\Phi_t)\psi^T(\Phi_t)$$

$$\frac{d}{dt}\varphi(t) = -\gamma\varphi(t) + \psi(\Phi_t)$$

$$\frac{d}{dt}b(t) = \varphi(t)c(\Phi_t)$$

$$\theta(t) := M(t)^{-1}b(t)$$

By law of large numbers,

$$\lim_{t \rightarrow \infty} \theta(t) = \theta^*$$

LSTD Learning

Discounted cost case

ODE formulation of the LSTD algorithm:

$$\begin{aligned}\frac{d}{dt}M(t) &= \psi(\Phi_t)\psi^T(\Phi_t) \\ \frac{d}{dt}\varphi(t) &= -\gamma\varphi(t) + \psi(\Phi_t) \\ \frac{d}{dt}b(t) &= \varphi(t)c(\Phi_t) \\ \theta(t) &:= M(t)^{-1}b(t)\end{aligned}$$

By law of large numbers,

$$\lim_{t \rightarrow \infty} \theta(t) = \theta^*$$

For average-cost ($\gamma = 0$) LSTD requires the existence of a regenerating state.

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

Idea: Approximate the gradient of h directly [R et al. 16, Devraj et al. 16]:

$$g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

Idea: Approximate the gradient of h directly [R et al. 16, Devraj et al. 16]:

$$g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Need to choose a function class \mathcal{H} for g (or ∇g)

- A finitely parameterized family of functions.
- A reproducing kernel Hilbert space (RKHS).

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

Idea: Approximate the gradient of h directly [R et al. 16, Devraj et al. 16]:

$$g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Need to choose a function class \mathcal{H} for g (or ∇g)

- A finitely parameterized family of functions.
- A reproducing kernel Hilbert space (RKHS).

Choice of basis is not an easy task

\implies RKHS framework is far easier to implement.

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \nabla \psi_i, \nabla \psi_j \rangle_{L^2}, \quad b_i = \langle \nabla \psi_i, \nabla h \rangle_{L^2}$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \nabla \psi_i, \nabla \psi_j \rangle_{L^2}, \quad b_i = \langle \nabla \psi_i, R_{U''} \nabla c \rangle_{L^2}$$

$$\text{Gen.resolvent kernel: } R_{U''} \nabla c(x) := \int_0^\infty \mathbb{E}_x \left[\exp \left(- \int_0^t U''(\Phi_s) \, ds \right) \nabla c(\Phi_t) \right] dt$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

$$\begin{aligned} M_{ij} &= \langle \nabla \psi_i, \nabla \psi_j \rangle_{L^2}, & b_i &= \langle \nabla \psi_i, R_{U''} \nabla c \rangle_{L^2} \\ & & &= \langle R_{U''}^\dagger \nabla \psi_i, \nabla c \rangle_{L^2} \end{aligned}$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

∇ -LSTD goal: $g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD-L: For Langevin diffusion, if $f, g \in L^2(\rho)$

$$\langle \nabla f, \nabla g \rangle_{L^2} = -\langle f, \mathcal{D}g \rangle_{L^2} = -\langle \mathcal{D}f, g \rangle_{L^2}.$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD-L: For Langevin diffusion, if $f, g \in L^2(\rho)$

$$\langle \nabla f, \nabla g \rangle_{L^2} = -\langle f, \mathcal{D}g \rangle_{L^2} = -\langle \mathcal{D}f, g \rangle_{L^2}.$$

Applying this and Poisson's equation $\mathcal{D}h = -\tilde{c}$:

$$\begin{aligned} \|\nabla h - \nabla g\|_{L^2}^2 &= \|\nabla h\|_{L^2}^2 + \|\nabla g\|_{L^2}^2 - 2\langle \nabla h, \nabla g \rangle_{L^2} \\ &= \|\nabla h\|_{L^2}^2 + \|\nabla g\|_{L^2}^2 + 2\langle \mathcal{D}h, g \rangle_{L^2} \end{aligned}$$

Differential LSTD Learning (∇ -LSTD)

Poisson's equation

$$\nabla\text{-LSTD goal: } g^* := \arg \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Challenge: the function h is not known,
and hence the objective function is not observable

∇ -LSTD-L: For Langevin diffusion, if $f, g \in L^2(\rho)$

$$\langle \nabla f, \nabla g \rangle_{L^2} = -\langle f, \mathcal{D}g \rangle_{L^2} = -\langle \mathcal{D}f, g \rangle_{L^2}.$$

Applying this and Poisson's equation $\mathcal{D}h = -\tilde{c}$:

$$\begin{aligned} \|\nabla h - \nabla g\|_{L^2}^2 &= \|\nabla h\|_{L^2}^2 + \|\nabla g\|_{L^2}^2 - 2\langle \nabla h, \nabla g \rangle_{L^2} \\ &= \|\nabla h\|_{L^2}^2 + \|\nabla g\|_{L^2}^2 - 2\langle \tilde{c}, g \rangle_{L^2} \end{aligned}$$

Differential LSTD Learning for Langevin (∇ -LSTD-L)

Poisson's equation

$$\nabla\text{-LSTD-L goal: } g^* := \arg \min_{g \in \mathcal{H}} \left\{ \|\nabla g\|_{L^2}^2 - 2\langle \tilde{c}, g \rangle_{L^2} \right\}$$

∇ -LSTD-L: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

Differential LSTD Learning for Langevin (∇ -LSTD-L)

Poisson's equation

$$\nabla\text{-LSTD-L goal: } g^* := \arg \min_{g \in \mathcal{H}} \left\{ \|\nabla g\|_{L^2}^2 - 2\langle \tilde{c}, g \rangle_{L^2} \right\}$$

∇ -LSTD-L: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

$$M_{ij} = \langle \nabla \psi_i, \nabla \psi_j \rangle_{L^2}$$

$$b_i = \langle \nabla \psi_i, \nabla h \rangle_{L^2} = \langle \psi_i, \tilde{c} \rangle_{L^2}$$

Differential LSTD Learning for Langevin (∇ -LSTD-L)

Poisson's equation

$$\nabla\text{-LSTD-L goal: } g^* := \arg \min_{g \in \mathcal{H}} \left\{ \|\nabla g\|_{L^2}^2 - 2\langle \tilde{c}, g \rangle_{L^2} \right\}$$

∇ -LSTD-L: For a linear parameterization

$$g = h^\theta := \sum_{i=1}^{\ell} \theta_i \psi_i \implies \nabla g = \sum_{i=1}^{\ell} \theta_i \nabla \psi_i$$

$$\theta^* = M^{-1}b$$

$$\begin{aligned} M_{ij} &= \langle \nabla \psi_i, \nabla \psi_j \rangle_{L^2} \\ &\approx \frac{1}{t} \int_0^t \nabla \psi(\Phi_s) \nabla \psi^T(\Phi_s) \, ds \end{aligned}$$

$$\begin{aligned} b_i &= \langle \nabla \psi_i, \nabla h \rangle_{L^2} = \langle \psi_i, \tilde{c} \rangle_{L^2} \\ &\approx \frac{1}{t} \int_0^t \psi_i(\Phi_s) \tilde{c}(\Phi_s) \, ds \end{aligned}$$

Differential LSTD Learning on RKHS (∇ -LSTD-RKHS)

Basics of RKHS

Choose a kernel function $K(x, y)$ that is

- *Symmetric*: $K(x, y) = K(y, x)$ for any $x, y \in \mathbb{R}^d$
- *Positive definite*: For any $\{x^i\} \subset \mathbb{R}^d$, matrix $\{M_{ij} := K(x^i, x^j)\}$ is positive definite.
- *Smooth*: $K \in C^2$

K defines a unique RKHS \mathcal{H} [Moore-Aronszajn theorem].

Inner product: If $g_\alpha = \sum_i \alpha_i K(x^i, \cdot)$ and $g_\beta = \sum_j \beta_j K(y^j, \cdot)$,

$$\langle g_\alpha, g_\beta \rangle_{\mathcal{H}} := \sum_{i,j} \alpha_i \beta_j K(x^i, y^j)$$

Reproducing property: $g_\alpha(x) = \langle g_\alpha, K(x, \cdot) \rangle_{\mathcal{H}}$, $x \in \mathbb{R}^d$.

Differential LSTD learning on RKHS

Empirical risk minimization (ERM)

Recall ∇ -LSTD-L goal:

$$g^* = \arg \min_{g \in \mathcal{H}} \left\{ \|\nabla g\|_{L^2}^2 - 2\langle \tilde{c}, g \rangle_{L^2} \right\}$$

Approximation via **empirical risk minimization (ERM)**:

$$\arg \min_{g \in \mathcal{H}} \underbrace{\frac{1}{N} \sum_{i=1}^N \left[\|\nabla g(x^i)\|^2 - 2\tilde{c}_N(x^i)g(x^i) \right]}_{\text{Empirical risk}} + \underbrace{\lambda \|g\|_{\mathcal{H}}^2}_{\text{Regularization}}$$

where $\tilde{c}_N(x) = c(x) - \frac{1}{N} \sum_{i=1}^N c(x^i)$, $x \in \mathbb{R}^d$

Differential LSTD Learning on RKHS (∇ -LSTD-RKHS) ∇ -LSTD-RKHS-Opt

Empirical risk minimization (ERM):

$$\arg \min_{g \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left[\|\nabla g(x^i)\|^2 - 2\tilde{c}_N(x^i)g(x^i) \right] + \lambda \|g\|_{\mathcal{H}}^2$$

Classical representer theorem [\[Wahba 70\]](#) is a remarkable result for ERMs in RKHS.

Differential LSTD Learning on RKHS (∇ -LSTD-RKHS)

∇ -LSTD-RKHS-Opt

Empirical risk minimization (ERM):

$$\arg \min_{g \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left[\|\nabla g(x^i)\|^2 - 2\tilde{c}_N(x^i)g(x^i) \right] + \lambda \|g\|_{\mathcal{H}}^2$$

Classical representer theorem [Wahba 70] is a remarkable result for ERM in RKHS.

Not applicable to our loss function due to gradient term.

Differential LSTD Learning on RKHS (∇ -LSTD-RKHS)

∇ -LSTD-RKHS-Opt

Empirical risk minimization (ERM):

$$\arg \min_{g \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left[\|\nabla g(x^i)\|^2 - 2\tilde{c}_N(x^i)g(x^i) \right] + \lambda \|g\|_{\mathcal{H}}^2$$

Extended Representer Theorem [Zhou 08]

If loss function $L(x, \cdot, \cdot)$ is convex on \mathbb{R}^{d+1} for each $x \in \mathbb{R}^d$, then the optimizer g^* over $g \in \mathcal{H}$ exists:

$$g^*(\cdot) = \sum_{i=1}^N \left[\beta_i^{0*} K(x^i, \cdot) + \sum_{k=1}^d \beta_i^{k*} \frac{\partial}{\partial x_k} K(x^i, \cdot) \right]$$

Differential LSTD Learning on RKHS (∇ -LSTD-RKHS)

∇ -LSTD-RKHS-N

Drawback : Complexity grows linearly with d , since $\beta^* \in \mathbb{R}^{(d+1) \times N}$

Simplified solution : By considering the finite-dimensional function class -
 $\mathcal{H}_N := \text{span}\{K_{x^i} : 1 \leq i \leq N\}$

$$g^*(y) = \sum_{j=1}^N \beta_j^* K(x^j, y)$$

$$\beta^* = M^{-1}b$$

Surprising empirical observation : Simplified solution does as good as the optimal solution for $d \leq 5$.

Applications to Nonlinear Filtering

Feedback Particle Filter

Goal: To obtain estimates of the state of a stochastic dynamical system based on noisy partial observations.

Applications to Nonlinear Filtering

Feedback Particle Filter

Goal: To obtain estimates of the state of a stochastic dynamical system based on noisy partial observations.

Kalman filter is optimal for a linear Gaussian system.

Applications to Nonlinear Filtering

Feedback Particle Filter

Goal: To obtain estimates of the state of a stochastic dynamical system based on noisy partial observations.

Kalman filter is optimal for a linear Gaussian system.

For nonlinear systems, conditional distribution fails to be Gaussian, cannot be captured by a finite set of parameters.

Applications to Nonlinear Filtering

Feedback Particle Filter

Goal: To obtain estimates of the state of a stochastic dynamical system based on noisy partial observations.

Kalman filter is optimal for a linear Gaussian system.

For nonlinear systems, conditional distribution fails to be Gaussian, cannot be captured by a finite set of parameters.

Particle filters are popular Monte-Carlo approximations of the nonlinear filter.

Applications to Nonlinear Filtering

Feedback Particle Filter

Problem:

Signal:
$$dX_t = a(X_t)dt + dB_t, \quad X_0 \sim \rho_0^*,$$

Observation:
$$dZ_t = c(X_t)dt + dW_t,$$

- $X_t \in \mathbb{R}^d$ is the state at time t .
- $\{Z_t : t \geq 0\}$ is the observation process.
- $a(\cdot), c(\cdot)$ are C^1 functions.
- $\{B_t\}, \{W_t\}$ are mutually independent Wiener processes.

Applications to Nonlinear Filtering

Feedback Particle Filter

Problem:

Signal:
$$dX_t = a(X_t)dt + dB_t, \quad X_0 \sim \rho_0^*,$$

Observation:
$$dZ_t = c(X_t)dt + dW_t,$$

- $X_t \in \mathbb{R}^d$ is the state at time t .
- $\{Z_t : t \geq 0\}$ is the observation process.
- $a(\cdot), c(\cdot)$ are C^1 functions.
- $\{B_t\}, \{W_t\}$ are mutually independent Wiener processes.
- $\rho_t^* := P(X_t | Z_s : s \leq t)$ is the posterior distribution.

Applications to Nonlinear Filtering

Feedback Particle Filter

Feedback particle filter (FPF) [\[Yang et al. 13\]](#) is motivated by techniques from mean-field optimal control.

Applications to Nonlinear Filtering

Feedback Particle Filter

Feedback particle filter (FPF) [Yang et al. 13] is motivated by techniques from mean-field optimal control.

N particles are propagated in the form of a controlled system.

$$dX_t^i = \underbrace{a(X_t^i)dt + dB_t^i}_{\text{Propagation}} + \underbrace{dU_t^i}_{\text{Update}}, \quad i = 1 \text{ to } N$$

- $X_t^i \in \mathbb{R}$ is the state of the i^{th} particle at time t
- U_t^i is the control input applied to i^{th} particle
- $\{B_t^i\}$ are mutually independent standard Wiener processes.

Applications to Nonlinear Filtering

Feedback Particle Filter

Feedback particle filter (FPF) [Yang et al. 13] is motivated by techniques from mean-field optimal control.

N particles are propagated in the form of a controlled system.

$$dX_t^i = \underbrace{a(X_t^i)dt + dB_t^i}_{\text{Propagation}} + \underbrace{dU_t^i}_{\text{Update}}, \quad i = 1 \text{ to } N$$

- $X_t^i \in \mathbb{R}$ is the state of the i^{th} particle at time t
- U_t^i is the control input applied to i^{th} particle
- $\{B_t^i\}$ are mutually independent standard Wiener processes.

Approximation of ρ_t^* :

$$\rho_t^* \approx \rho_t^{(N)}(A) = \frac{1}{N} \sum_{i=1}^N \mathbb{I}\{X_t^i \in A\}, \quad A \subset \mathbb{R}.$$

Applications to Nonlinear Filtering

Feedback Particle Filter

$$dU_t^i = \mathbf{K}_t(\mathbf{X}_t^i) \circ \overbrace{(dZ_t - \frac{1}{2}[c(\mathbf{X}_t^i) + \hat{c}_t]dt)}^{dI_t^i}$$

I_t^i : Innovations process.

\mathbf{K}_t : FPF gain, similar in nature to the Kalman gain.


Applications to Nonlinear Filtering

Feedback Particle Filter

$$dU_t^i = \mathbf{K}_t(X_t^i) \circ \overbrace{(dZ_t - \frac{1}{2}[c(X_t^i) + \hat{c}_t]dt)}^{dI_t^i}$$

I_t^i : Innovations process.

K_t : FPF gain, similar in nature to the Kalman gain.




Finite- N implementation

Applications to Nonlinear Filtering

Feedback Particle Filter

KF:
$$d\mu_t = a(\mu_t)dt + \underbrace{K_{\text{kal}}(dZ_t - c(\mu_t)dt)}_{\text{update}}$$




Kalman filter

Applications to Nonlinear Filtering


Feedback Particle Filter

KF:
$$d\mu_t = a(\mu_t)dt + \underbrace{K_{\text{kal}}(dZ_t - c(\mu_t)dt)}_{\text{update}}$$

FPF:
$$dX_t^i = a(X_t^i)dt + dB_t^i + \underbrace{K_t(X_t^i) \circ (dZ_t - \frac{1}{2}[c(X_t^i) + \hat{c}_t]dt)}_{\text{update}}$$



Kalman filter



Feedback particle filter (FPF)

Applications to Nonlinear Filtering

FPF gain function

Representation: $K_t = \nabla h$

h solves Poisson's equation: $\mathcal{D}h = -\nabla U \cdot \nabla h + \Delta h = -\tilde{c}$.

Applications to Nonlinear Filtering

FPF gain function

Representation: $K_t = \nabla h$

h solves **Poisson's equation**: $\mathcal{D}h = -\nabla U \cdot \nabla h + \Delta h = -\tilde{c}$.

Approximations to K can be obtained by

$$\min_{g \in \mathcal{H}} \|K - \hat{K}\|_{L^2}^2 = \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Applications to Nonlinear Filtering

FPF gain function

Representation: $K_t = \nabla h$

h solves **Poisson's equation**: $\mathcal{D}h = -\nabla U \cdot \nabla h + \Delta h = -\tilde{c}$.

Approximations to K can be obtained by

$$\min_{g \in \mathcal{H}} \|K - \hat{K}\|_{L^2}^2 = \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

Can be solved using ∇ -LSTD learning

Applications to Nonlinear Filtering

FPF gain function

Representation: $K_t = \nabla h$

h solves **Poisson's equation**: $\mathcal{D}h = -\nabla U \cdot \nabla h + \Delta h = -\tilde{c}$.

Approximations to K can be obtained by

$$\min_{g \in \mathcal{H}} \|K - \hat{K}\|_{L^2}^2 = \min_{g \in \mathcal{H}} \|\nabla h - \nabla g\|_{L^2}^2$$

FPF implementation requires online gain estimation for each t .

- ∇ -LSTD-RKHS with optimal mean
- ∇ -LSTD-RKHS with memory

Applications to Nonlinear filtering

▽-LSTD-RKHS-OM

Constant gain approximation for K is the minimizer obtained over all deterministic vectors:

$$\hat{K}^* := \arg \min_{\hat{K} \in \mathbb{R}^d} \|K - \hat{K}\|_{L^2}^2$$

Applications to Nonlinear filtering

▽-LSTD-RKHS-OM

Constant gain approximation for K is the minimizer obtained over all deterministic vectors:

$$\hat{K}^* := \arg \min_{\hat{K} \in \mathbb{R}^d} \|K - \hat{K}\|_{L^2}^2$$

Solution is evidently the mean, $\hat{K}^* = E[K]$.

$$\hat{K}_k^* = \langle K, e_k \rangle_{L^2}$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Constant gain approximation for K is the minimizer obtained over all deterministic vectors:

$$\hat{K}^* := \arg \min_{\hat{K} \in \mathbb{R}^d} \|K - \hat{K}\|_{L^2}^2$$

Solution is evidently the mean, $\hat{K}^* = E[K]$.

$$\begin{aligned} \hat{K}_k^* &= \langle K, e_k \rangle_{L^2} \\ &= \langle \nabla h, \nabla x_k \rangle_{L^2} \end{aligned}$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Constant gain approximation for K is the minimizer obtained over all deterministic vectors:

$$\hat{K}^* := \arg \min_{\hat{K} \in \mathbb{R}^d} \|K - \hat{K}\|_{L^2}^2$$

Solution is evidently the mean, $\hat{K}^* = E[K]$.

$$\begin{aligned} \hat{K}_k^* &= \langle K, e_k \rangle_{L^2} \\ &= \langle \nabla h, \nabla x_k \rangle_{L^2} \\ &= -\langle \mathcal{D}h, x_k \rangle_{L^2} = \langle \tilde{c}, x_k \rangle_{L^2} \end{aligned}$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Constant gain approximation for K is the minimizer obtained over all deterministic vectors:

$$\hat{K}^* := \arg \min_{\hat{K} \in \mathbb{R}^d} \|K - \hat{K}\|_{L^2}^2$$

Solution is evidently the mean, $\hat{K}^* = E[K]$.

$$\begin{aligned} \hat{K}_k^* &= \langle K, e_k \rangle_{L^2} \\ &= \langle \nabla h, \nabla x_k \rangle_{L^2} \\ &= -\langle \mathcal{D}h, x_k \rangle_{L^2} = \langle \tilde{c}, x_k \rangle_{L^2} \end{aligned}$$

Empirical approximation:

$$\hat{K}_k^* \approx \frac{1}{N} \sum_{i=1}^N [c(x^i) - \hat{c}] x_k^i$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Redefine the approximation to K as,

$$\nabla g = \hat{K}^* + \nabla \tilde{g}$$

Modified ERM with constraints is:

$$\begin{aligned} \tilde{g}^* &:= \arg \min_{\tilde{g} \in \mathcal{H}} \|\nabla h - \hat{K}^* - \nabla \tilde{g}\|_{L_2}^2 \\ \text{s.t.} \quad &\langle \partial_{x_k} \tilde{g}, 1 \rangle_{L_2} = 0, \quad 1 \leq k \leq d \end{aligned}$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Redefine the approximation to K as,

$$\nabla g = \widehat{K}^* + \nabla \tilde{g}$$

Modified ERM with constraints is:

$$\begin{aligned} \tilde{g}^* &:= \arg \min_{\tilde{g} \in \mathcal{H}} \|\nabla h - \widehat{K}^* - \nabla \tilde{g}\|_{L_2}^2 \\ \text{s.t. } &\langle \partial_{x_k} \tilde{g}, 1 \rangle_{L_2} = 0, \quad 1 \leq k \leq d \end{aligned}$$

Solution obtained by finding a saddle point for the Lagrangian

$$L(\tilde{g}, \mu) := \|\nabla h - \widehat{K}^* - \nabla \tilde{g}\|_{L_2}^2 + \langle \mu, \nabla \tilde{g} \rangle_{L_2}$$

where $\mu \in \mathbb{R}^d$ are the Lagrange multipliers.

Applications to Nonlinear filtering

∇ -LSTD-RKHS-OM

Redefine the approximation to K as,

$$\nabla g = \widehat{K}^* + \nabla \tilde{g}$$

Modified ERM with constraints is:

$$\begin{aligned} \tilde{g}^* &:= \arg \min_{\tilde{g} \in \mathcal{H}} \|\nabla h - \widehat{K}^* - \nabla \tilde{g}\|_{L_2}^2 \\ \text{s.t.} \quad &\langle \partial_{x_k} \tilde{g}, 1 \rangle_{L_2} = 0, \quad 1 \leq k \leq d \end{aligned}$$

Using $\mathcal{H}_N := \text{span}\{K_{x^j} : 1 \leq j \leq N\}$, β and μ can be obtained by solving $N + d$ linear equations

$$K = \widehat{K}^* + \nabla \tilde{g}^*$$

Applications to Nonlinear filtering

∇ -LSTD-RKHS-memory

Gain updates are done at $t = n\delta$, where δ is the inter-sampling time.

Continuity: $K_n = K_{t_n} \approx K_{t_{n-1}}$ if $\delta \approx 0$.

Applications to Nonlinear filtering

∇ -LSTD-RKHS-memory

Gain updates are done at $t = n\delta$, where δ is the inter-sampling time.

Continuity: $K_n = K_{t_n} \approx K_{t_{n-1}}$ if $\delta \approx 0$.

Adding a regularization term to the loss function:

$$g_n^* := \arg \min_{g \in \mathcal{H}} \frac{1}{N} \sum_{j=1}^N L_n(x_n^j, g, \nabla g) + \lambda \|g\|_{\mathcal{H}}^2$$

$$L_n(x, g, \nabla g) := \|\nabla g(x)\|^2 - 2\tilde{c}_N(x)g(x) + \underbrace{\lambda_{mem} \|\nabla g(x) - \nabla g_{n-1}(x)\|^2}_{\text{continuity penalty}}$$

$$\beta_n^* = M^{-1}b$$

Application to nonlinear filtering


Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $T = 10^4, 10^5, 10^6$ with $\delta = 0.01$, $N = 1000$

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $T = 10^4, 10^5, 10^6$ with $\delta = 0.01$, $N = 1000$




∇ -LSTD with $\psi_i = x^i \rho_1(x)$, $\psi_{i+1} = x^i \rho_2(x)$ with $1 \leq i \leq 5$.

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




∇ -LSTD-L with $\psi_i = x^i \rho_1(x)$, $\psi_{i+1} = x^i \rho_2(x)$ with $1 \leq i \leq 5$.

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




∇ -LSTD-RKHS-N with Gaussian kernel, $\varepsilon = 0.1$ and $\lambda = 10^{-2}$ (best)

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




∇ -LSTD-RKHS-2N with Gaussian kernel, $\varepsilon = 0.1$ and $\lambda = 10^{-2}$ (best)

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




∇ -LSTD-RKHS-OM with Gaussian kernel, $\varepsilon = 0.1$ and $\lambda = 10^{-2}$ (best)

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




Markov kernel approximation [Taghvaei et al. 18] with $\epsilon = 0.1$ (best)

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: For a fixed t , ρ_t a Gaussian mixture
 $c(x) \equiv x$, $d = 1$
 $N = 100, 500, 1000$




Histogram of MSEs obtained from 1000 independent trials

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: $\rho(x) = \prod_{k=1}^d \rho_k(x_k)$, each ρ_k a Gaussian mixture
 $c(x) = C^T x$, where $C = \mathbb{I}_d$
 $d = 2, 5, 10, \quad N = 1000$






∇ -LSTD-RKHS-OM

Application to nonlinear filtering

Numerical example - Gain approximation for a fixed t

Example: $\rho(x) = \prod_{k=1}^d \rho_k(x_k)$, each ρ_k a Gaussian mixture
 $c(x) = C^T x$, where $C = \mathbb{I}_d$
 $d = 2, 5, 10, \quad N = 1000$




Markov kernel approximation

Application to nonlinear filtering

Numerical example - Gain for a nonlinear oscillator model

Example: $\rho(x)$ is a mixture of von Mises densities on a circle
 $d\vartheta = \omega dt + \sigma_B dB_t \mod 2\pi$,
 $dZ_t = c(\vartheta)dt + \sigma_W dW_t$, $c(\vartheta) = \frac{1}{2}[1 + \cos(\vartheta)]$




∇ -LSTD-L with $\psi_i = \sin(ix)$, $\psi_{i+1} = \cos(ix)$ with $1 \leq i \leq \ell/2$, $\ell = 4, 6, 8$.

Applications to Nonlinear filtering

Numerical example - Parameter estimation


Example: Parameter estimation with bimodal prior
Observations: parameter plus additive noise with $\sigma_W = 1$.



Applications to Nonlinear filtering

Numerical example - Parameter estimation

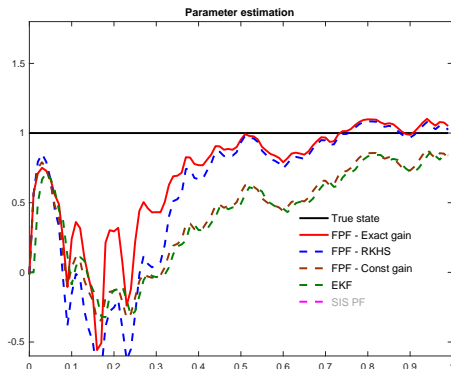
Example: Parameter estimation with bimodal prior
Observations: parameter plus additive noise with $\sigma_W = 1$.



Applications to Nonlinear filtering

Numerical example - Parameter estimation


Example: Parameter estimation with bimodal prior
Observations: parameter plus additive noise with $\sigma_W = 1$.



Applications to Nonlinear filtering

Numerical example - Parameter estimation

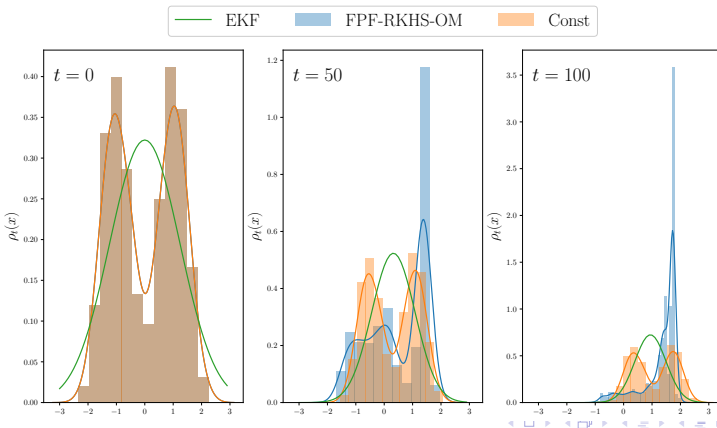
Example: Parameter estimation with bimodal prior
 Observations: parameter plus additive noise with $\sigma_W = 1$.



Applications to Nonlinear filtering

Numerical example - Parameter estimation


Example: Parameter estimation with bimodal prior
 Observations: parameter plus additive noise with $\sigma_W = 1$.



Applications to Nonlinear filtering

Numerical example - Nonlinear 2d ship dynamics model

Example: Nonlinear ship dynamics model in 2d.
 Observations: $c(x) = \arctan(x_1/x_2)$ with std. deviation $\approx 18^\circ$.



Filter	Σ_1	Lost track
RKHS-OM	0.90	4
RKHS-mem.	0.91	7
Const.	1.30	14
SIR PF	3.14	57
EKF	6.52	93

Applications to MCMC

Introduction to MCMC

In many applications, we need to compute

$$\eta = \int c(x) \rho(x) \, dx$$

- $c: \mathbb{R}^\ell \rightarrow \mathbb{R}$ is a measurable function.
- ρ is a target probability density in \mathbb{R}^ℓ .

Markov-Chain Monte Carlo (MCMC) methods provide numerical algorithms to obtain estimates:

$$\eta_t = \frac{1}{t} \int_0^t c(\Phi(s)) \, ds$$

Φ is a Markov process with steady state distribution ρ .

Applications to MCMC

Asymptotic Variance

Estimates η_t obey a Central Limit Theorem,

$$\sqrt{t}(\eta_t - \eta) \xrightarrow{d} N(0, \gamma^2)$$

Representation in terms of h [see eg: Glynn & Meyn 96]:

$$\gamma^2 = 2\langle h, \tilde{c} \rangle$$

Applications to MCMC

Asymptotic Variance

Estimates η_t obey a Central Limit Theorem,

$$\sqrt{t}(\eta_t - \eta) \xrightarrow{d} N(0, \gamma^2)$$

Representation in terms of h [see eg: Glynn & Meyn 96]:

$$\begin{aligned}\gamma^2 &= 2\langle h, \tilde{c} \rangle \\ &= 2\|\nabla h\|_{L^2}^2 \\ &\text{(For Langevin diffusion)}\end{aligned}$$

Applications to MCMC

Control Variates

Goal: To minimize asymptotic variance.

Idea: Modify the estimator using control variates [Henderson 97, CTCN¹]

$$c_g = c + \underbrace{\mathcal{D}g}_{\text{Control variate}}, \quad \text{where } g \in \mathcal{H}$$

$$\eta_t^g = \frac{1}{t} \int_0^t c_g(\Phi_s) ds$$

For any $g \in C^2 \cap L^2(\rho)$, $\langle \mathcal{D}g, 1 \rangle_{L^2} = 0$.

¹Control Techniques for Complex Networks, S.Meyn

Applications to MCMC

Optimal control variates

$$\mathcal{D}(h - g) = -c_g + \eta$$

$h - g$ is the solution to Poisson's equation with \tilde{c}_g as the forcing function.

Applications to MCMC

Optimal control variates

$$\mathcal{D}(h - g) = -c_g + \eta$$

$h - g$ is the solution to Poisson's equation with \tilde{c}_g as the forcing function.

Asymptotic variance of the new estimator:

$$\gamma_g^2 = 2\langle h - g, \tilde{c}_g \rangle_{L^2}, \quad \tilde{c}_g = c_g - \eta$$

Applications to MCMC

Optimal control variates

$$\mathcal{D}(h - g) = -c_g + \eta$$

$h - g$ is the solution to Poisson's equation with \tilde{c}_g as the forcing function.

Asymptotic variance of the new estimator:

$$\begin{aligned}\gamma_g^2 &= 2\langle h - g, \tilde{c}_g \rangle_{L^2}, \quad \tilde{c}_g = c_g - \eta \\ &= 2\|\nabla h - \nabla g\|_{L^2}^2, \quad (\text{For Langevin diffusion})\end{aligned}$$

Applications to MCMC

Optimal control variates

$$\mathcal{D}(h - g) = -c_g + \eta$$

$h - g$ is the solution to Poisson's equation with \tilde{c}_g as the forcing function.

Asymptotic variance of the new estimator:

$$\begin{aligned}\gamma_g^2 &= 2\langle h - g, \tilde{c}_g \rangle_{L^2}, \quad \tilde{c}_g = c_g - \eta \\ &= 2\|\nabla h - \nabla g\|_{L^2}^2, \quad (\text{For Langevin diffusion})\end{aligned}$$

Can be minimized using ∇ -LSTD algorithms.

Applications to MCMC

Numerical Examples - Variance v Asymptotic variance

Variance v Asymptotic variance

$$\sigma^2 = \langle \tilde{c}, \tilde{c} \rangle_{L^2} = R(0)$$

- Variance

Applications to MCMC

Numerical Examples - Variance v Asymptotic variance

Variance v Asymptotic variance

$$\sigma^2 = \langle \tilde{c}, \tilde{c} \rangle_{L^2} = R(0) \quad - \text{Variance}$$

$$\gamma^2 = 2 \langle h, \tilde{c} \rangle_{L^2} = \int_{-\infty}^{\infty} R(s) ds \quad - \text{Asymptotic variance}$$

Applications to MCMC

Numerical Examples - Variance v Asymptotic variance

Variance v Asymptotic variance

$$\sigma^2 = \langle \tilde{c}, \tilde{c} \rangle_{L^2} = R(0) \quad - \text{ Variance}$$

$$\gamma^2 = 2 \langle h, \tilde{c} \rangle_{L^2} = \int_{-\infty}^{\infty} R(s) ds \quad - \text{ Asymptotic variance}$$

Minimizing σ^2 is easier than minimizing γ^2 (ZV method [\[Papamarkou 14\]](#))

But does minimizing $\sigma^2 \implies$ minimizing γ^2 ?

Applications to MCMC

Numerical Examples - Variance v Asymptotic variance

Variance v Asymptotic variance

$$\sigma^2 = \langle \tilde{c}, \tilde{c} \rangle_{L^2} = R(0) \quad - \text{Variance}$$

$$\gamma^2 = 2 \langle h, \tilde{c} \rangle_{L^2} = \int_{-\infty}^{\infty} R(s) ds \quad - \text{Asymptotic variance}$$

Minimizing σ^2 is easier than minimizing γ^2 (ZV method [\[Papamarkou 14\]](#))



But does minimizing $\sigma^2 \implies$ minimizing γ^2 ? **NO !**

Applications to MCMC

Numerical Examples - Variance vs Asymptotic variance

Example: Unadjusted Langevin algorithm (ULA)



$$c(x) \equiv x$$



Applications to MCMC

Numerical Examples - Variance vs Asymptotic variance

Example: Unadjusted Langevin algorithm (ULA)
 $c(x) \equiv x$

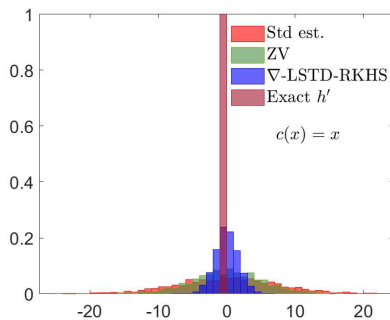
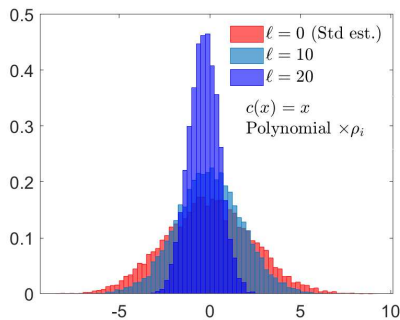


Applications to MCMC

Numerical Examples - ULA

Example: Unadjusted Langevin algorithm (ULA)

$$c(x) \equiv x, \rho \sim 0.5\mathcal{N}(-1, 0.4472) + 0.5\mathcal{N}(1, 0.4472)$$





Applications to MCMC

Numerical Examples - RWM

Example: Random walk Metropolis (RWM)

$$c(x) \equiv x, \rho \sim 0.5\mathcal{N}(-1, 0.4472) + 0.5\mathcal{N}(1, 0.4472)$$



Theoretical justification based on [\[Brosse et al. 19\]](#)



Applications to MCMC

Numerical Examples - Logistic Regression with RWM sampling

Example: Logistic Regression for Swiss bank notes
RWM sampling



Box plots of estimates of Θ_4 .

Conclusions

- Differential LSTD learning algorithms to approximate solution to Poisson's equation for the Langevin diffusion.
 - RKHS based approaches solve the basis selection problem and enable easy extensions to higher dimensions.
 - Optimal mean gain is a useful design tool and makes the algorithm more robust to ε and λ .
- Two interesting applications
 - Gain function approximation in feedback particle filter.
 - Asymptotic variance reduction in MCMC algorithms.
- Recent research extended to include reversible Markov chains.

Future Work

Analysis:

- Error analysis of the ∇ -LSTD-RKHS method, which could lead to proper choices of hyper parameters.
- A more thorough comparison of the various gain approximation algorithms.
- **Nagging question:** Why reduced complexity solution is as good as the optimal?

Future Work

Analysis:

- Error analysis of the ∇ -LSTD-RKHS method, which could lead to proper choices of hyper parameters.
- A more thorough comparison of the various gain approximation algorithms.
- **Nagging question:** Why reduced complexity solution is as good as the optimal?

Algorithm:

- ∇ -LSTD-RKHS algorithm with a differential regularizer. This is limited by the scope of representer theorem.
- An algorithm based on semiparametric representer theorem, that allows the use of additional parameterized functions.

Future Work

Analysis:

- Error analysis of the ∇ -LSTD-RKHS method, which could lead to proper choices of hyper parameters.
- A more thorough comparison of the various gain approximation algorithms.
- **Nagging question:** Why reduced complexity solution is as good as the optimal?

Algorithm:

- ∇ -LSTD-RKHS algorithm with a differential regularizer. This is limited by the scope of representer theorem.
- An algorithm based on semiparametric representer theorem, that allows the use of additional parameterized functions.

Applications:

- Real time filtering problem - Potential application to battery SOC estimation is being explored currently.
- Application of FPF to control in the context of POMDPs.

References



A. Radhakrishnan, A. Devraj and S. Meyn, "Learning techniques for feedback particle filter design," *2016 IEEE 55th Conference on Decision and Control (CDC)*, Las Vegas, NV, 2016.



A. Radhakrishnan, S. Meyn, "Feedback particle filter design using a differential-loss reproducing kernel Hilbert space," *2018 American Control Conference (ACC)*, Milwaukee, WI, 2018.



A. Radhakrishnan, S. Meyn, "Gain function tracking in the feedback particle filter," *2019 American Control Conference (ACC)*, Philadelphia, PA, 2019.



S.P.Meyn, "Control Techniques for Complex Networks", *Cambridge University Press*, Dec 2007.



S. Henderson. *Variance Reduction Via an Approximating Markov Process*. PhD thesis, Stanford University, Stanford, California, 1997.



T. Yang, P. G. Mehta and S. P. Meyn, "Feedback Particle Filter," in *IEEE Transactions on Automatic Control*, vol. 58, no. 10, pp. 2465-2480, Oct. 2013.



A. M. Devraj and S. P. Meyn, "Differential LSTD learning for value function approximation," *2016 IEEE 55th Conference on Decision and Control (CDC)*, Las Vegas, NV, 2016.



D.X. Zhou, "Derivative reproducing properties for kernel methods in learning theory," *Journal of Computational and Applied Mathematics*, Vol. 220, Issues 1?2, 2008.

Thank You!

Questions?